

Deploying R for production

2019-01-23

Stockholm R User Group



Holger Hellebro
IT Architect & Data Scientist
Cognitive & Analytics
IBM Global Business Services

<https://www.linkedin.com/in/holgerhellebro>
Twitter: @holken



Can we deploy R code?

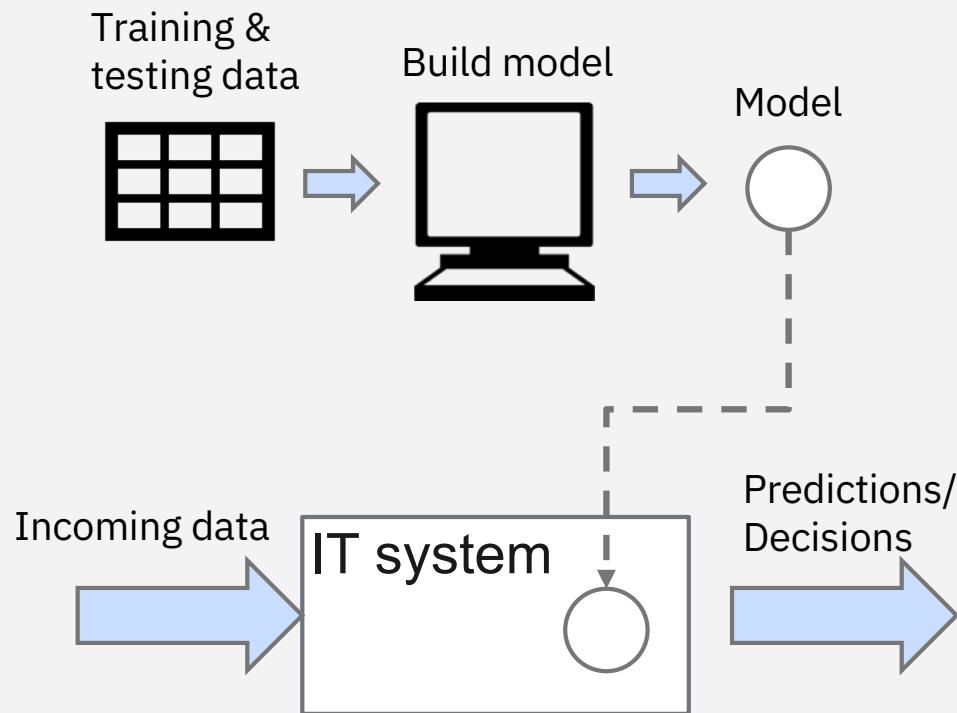
Why would we?

Will it scale?

How?

Security?

Why deploy R stuff?



Models for scoring
Visualizations
Data preparation
Feature extraction
Verify with more data

Deploy or rewrite?

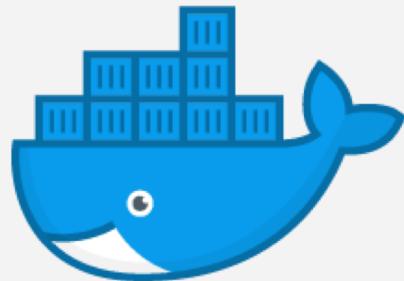
Why deploy?

- Reuse code
- Consistency
- Easier to maintain

Why rewrite?

- Minimize number of languages
- IT team skills
- Extreme performance / latency

Modern cloud technologies



docker

Helps us package code and data.



kubernetes

Helps us run it on a cluster.

Docker images & containers

Means of packaging and running software

Independent of host system

Include all dependencies

Light weight



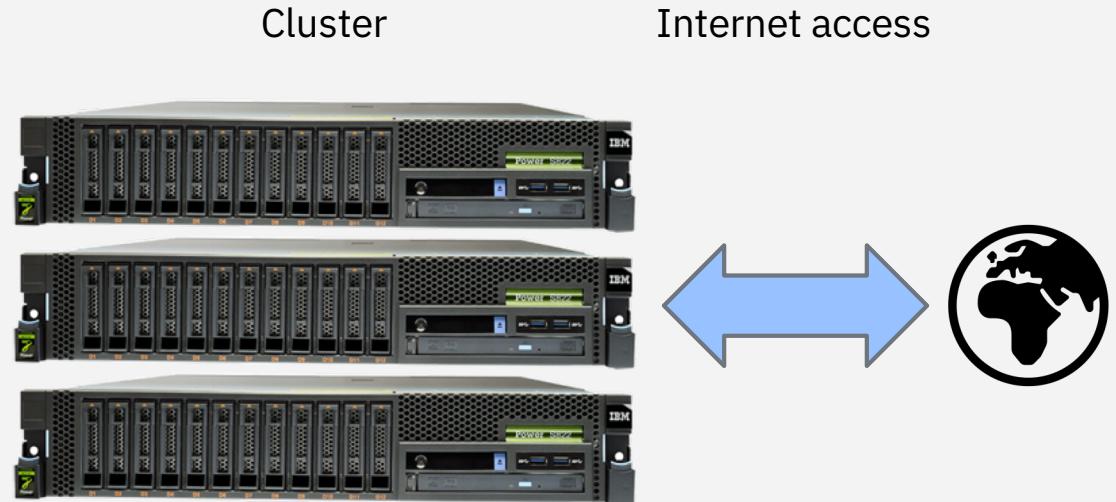
“Dockerfile” tells docker how to build into an image

Images contain all dependencies, but not a full Operating System

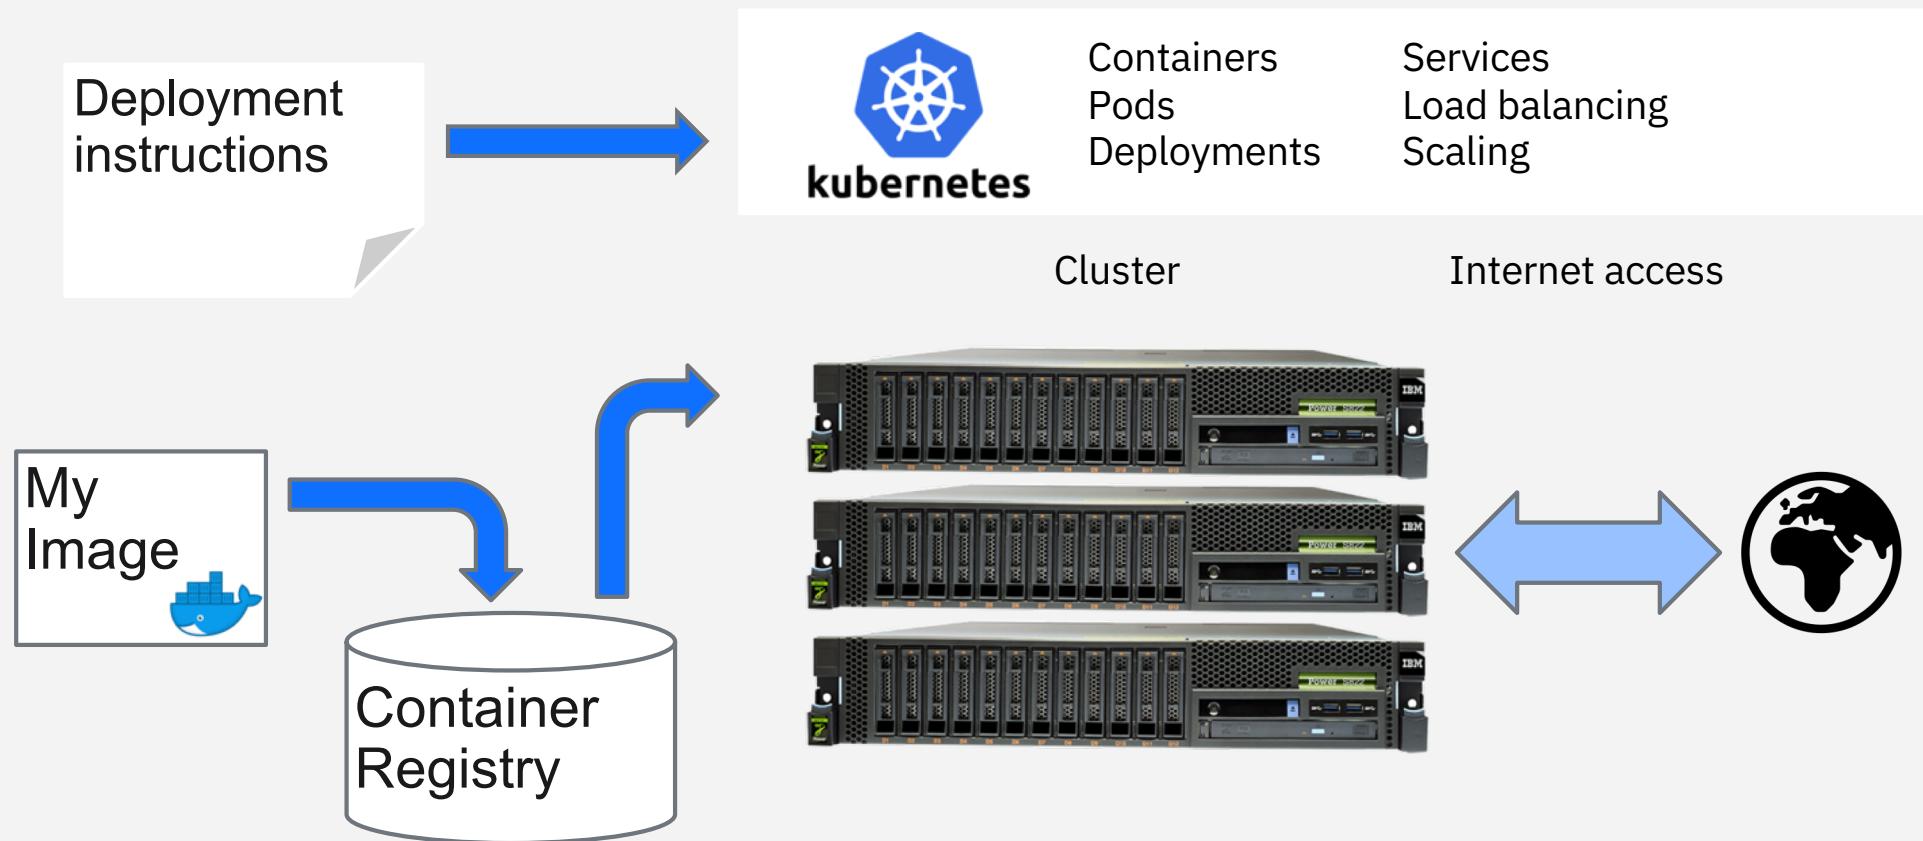
Container images are uploaded to a **registry** to simplify deployment



Kubernetes



Kubernetes



Deploying a Shiny app



Developing and building

app.R

```
library(shiny)

ui <- fluidPage(
  titlePanel("My Shiny App"),
  mainPanel(
    actionButton("button", "Randomize"),
    plotOutput("distPlot")
  )
)

server <- function(input, output) {
  output$distPlot <- renderPlot({
    input$button
    hist(rnorm(1000, 0, 1), breaks = 30,
         col = 'darkgray', border = 'white')
  })
}

shinyApp(ui = ui, server = server)
```

Dockerfile



```
FROM rocker/shiny:latest
WORKDIR /srv/shiny-server/
COPY . .
EXPOSE 3838
CMD ["/usr/bin/shiny-server.sh"]
```

```
$ docker build . -t holken/myshiny
$ docker run -p 3838:3838 holken/myshiny
$ docker push holken/myshiny
```

Deploy

Deployment descriptor

- What container(s)
 - Number of instances (replicas)
 - Secrets
 - Persistent storage
 - Resource requirements & limits
- ...

deployment.yaml



kubernetes

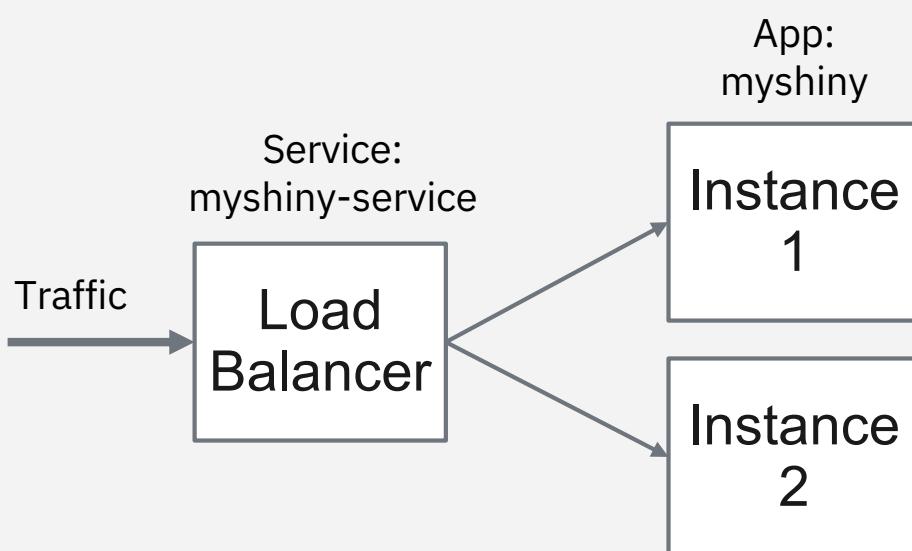
```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: myshiny-deployment
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: myshiny
    spec:
      containers:
      - name: myshiny
        image:
          docker.io/holken/myshiny:latest
```

Networking



Isolated by default

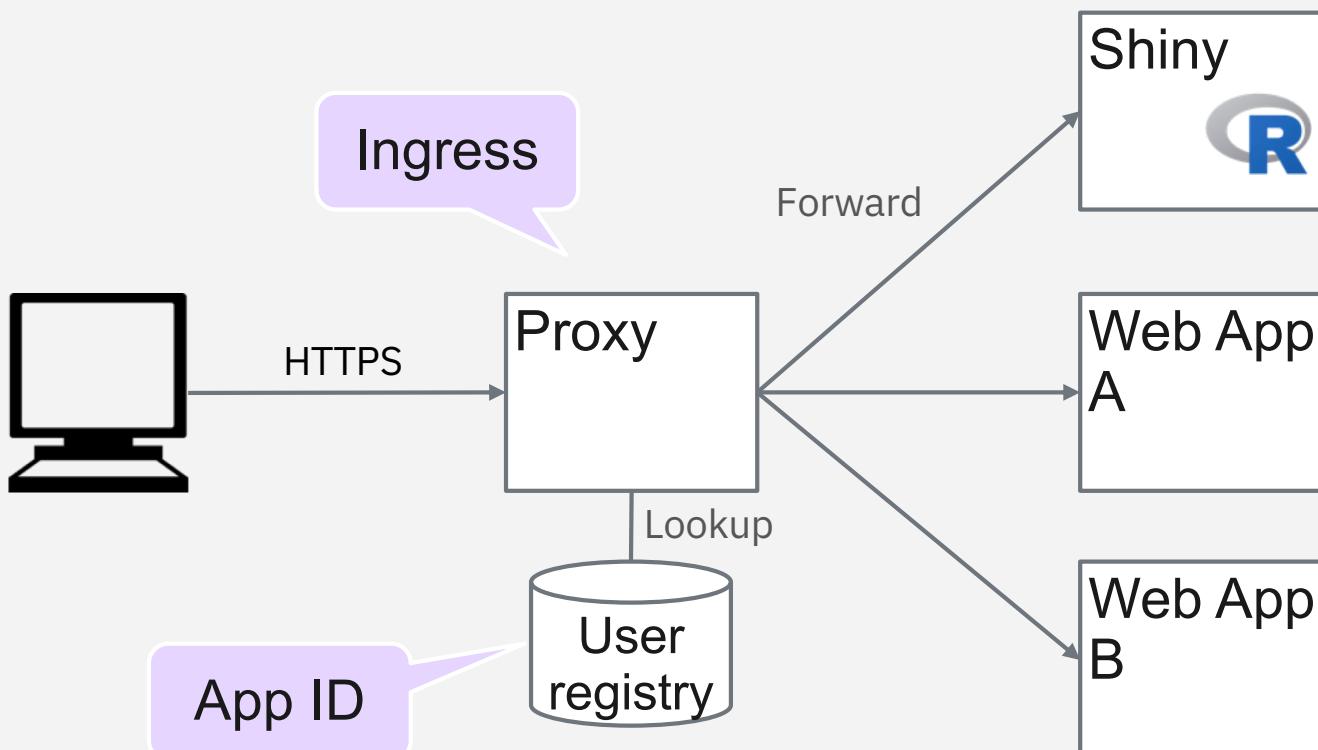
“Service” defines networking



service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: myshiny-service
spec:
  ports:
  - port: 80
    targetPort: 3838
    protocol: TCP
  type: LoadBalancer
  selector:
    app: myshiny
```

Making it secure with Ingress and App ID



[IBM Cloud Docs](#) / [IBM Cloud Kubernetes Service](#)

Exposing apps with Ingress

Last Updated: 2019-01-21 | [Edit in GitHub](#)



App ID
Lite • IBM

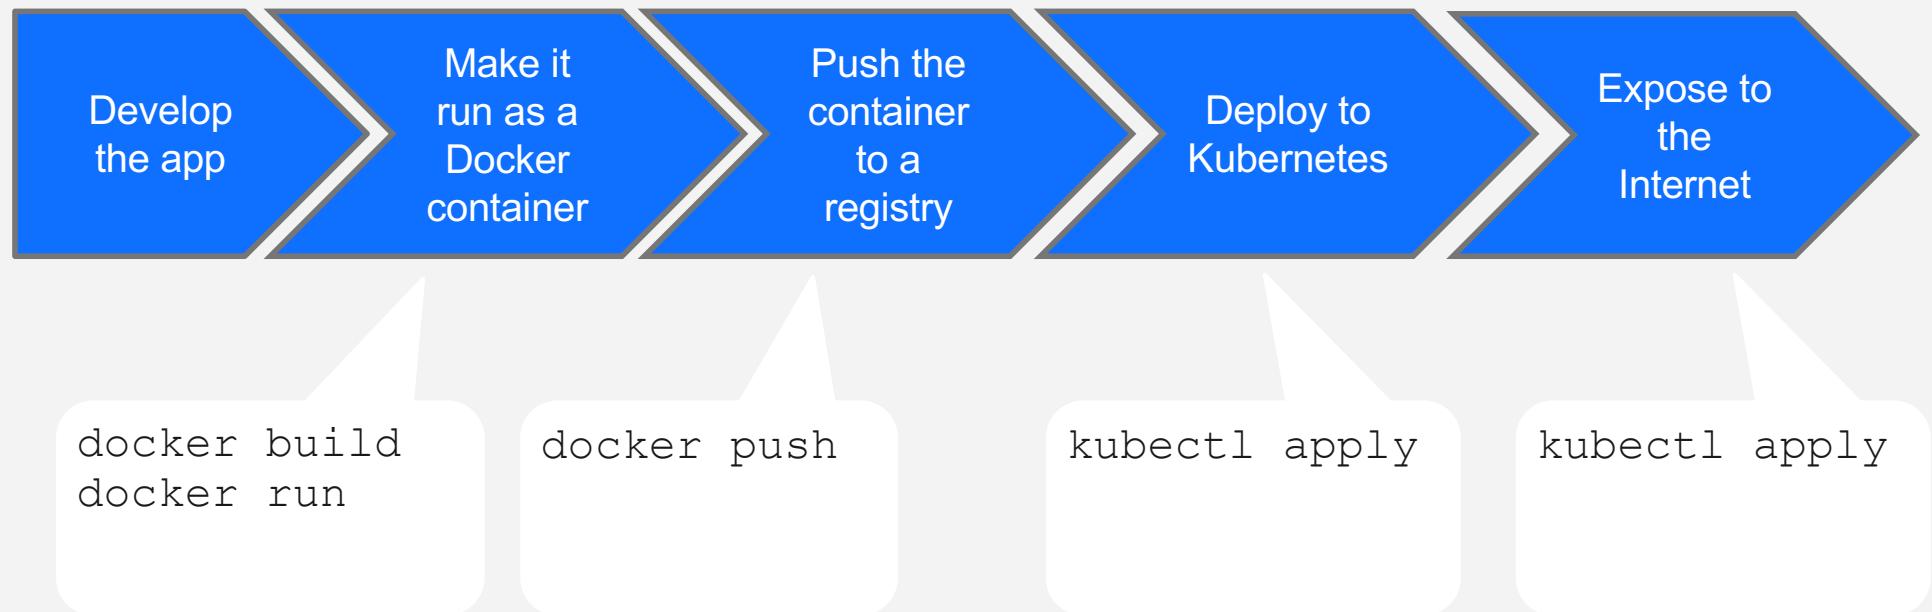
User Authentication and User Profiles for your apps.

https://console.bluemix.net/docs/containers/cs_ingress.html#ingress
https://console.bluemix.net/docs/containers/cs_annotations.html#appid-auth

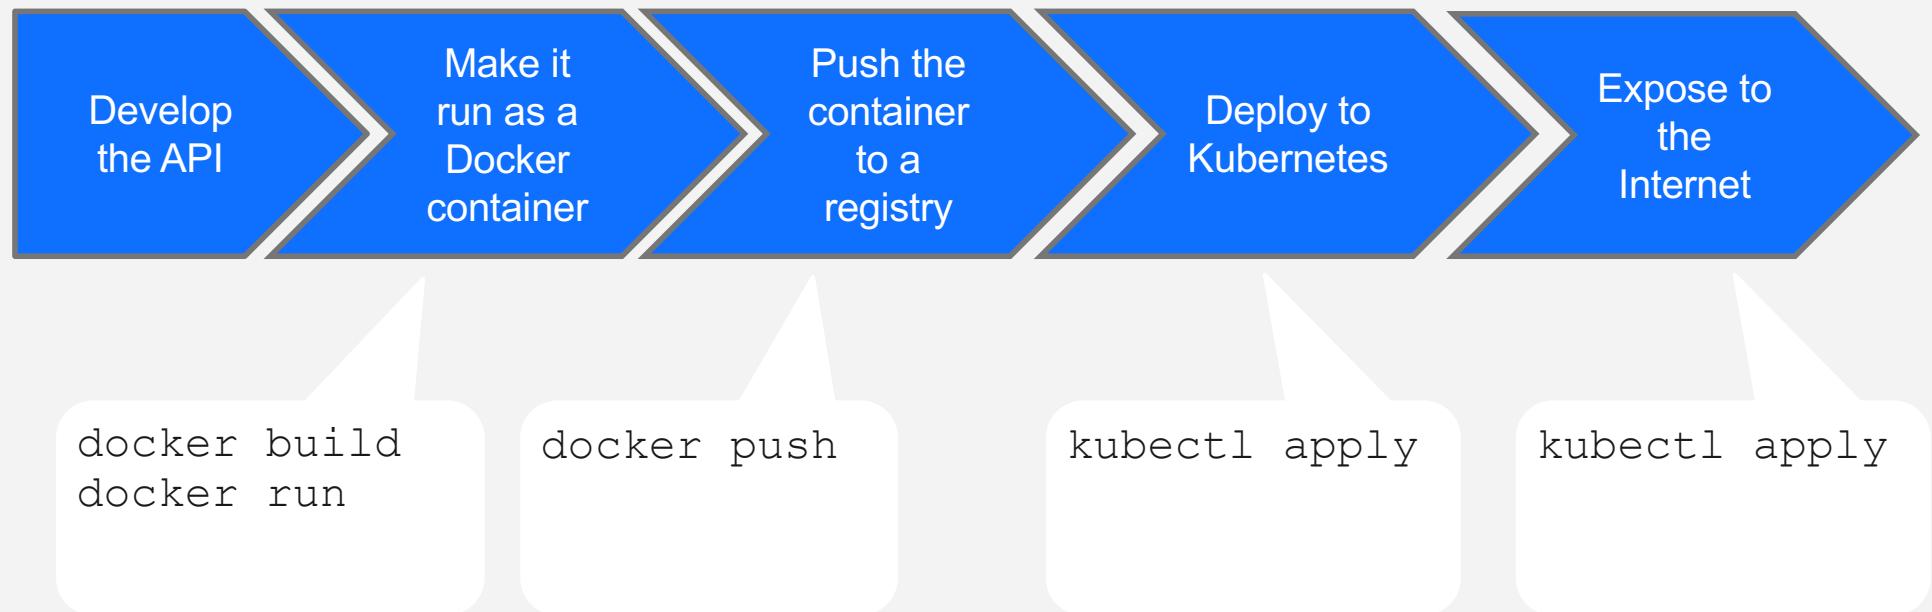
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-test-ingress
  annotations:
    ingress.bluemix.net/rewrite-path: "serviceName=myshiny-service rewrite=/"
    ingress.bluemix.net/appid-auth: "bindSecret=binding-appid-test namespace=default
requestType=web serviceName=myshiny-service,web-service"
spec:
  tls:
    - hosts:
        - growsmarter-test.eu-de.containers.appdomain.cloud
      secretName: growsmarter-test
  rules:
    - host: growsmarter-test.eu-de.containers.appdomain.cloud
      http:
        paths:
          - path: /shiny/
            backend:
              serviceName: myshiny-service
              servicePort: 3838
          - path: /
            backend:
              serviceName: web-service
              servicePort: 80
```



Deploying a Shiny app



Deploying an API



Building APIs with plumber

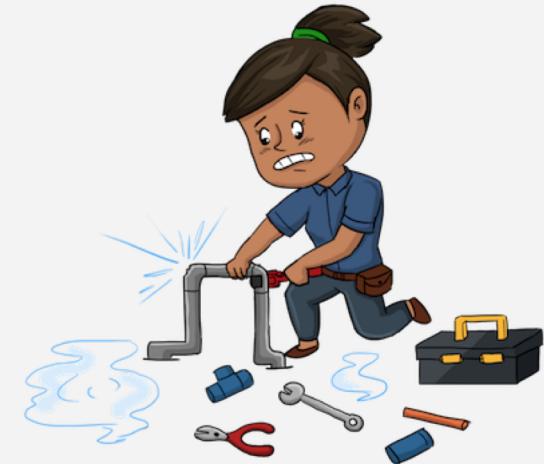
Step 1. Building the model

create-model.R

```
library(ggplot2) ## for dataset
library(readr) ## for write_rds
data(diamonds)

model <- lm(price ~ carat, data = diamonds)

write_rds(model, "lm-model.rds")
```



<https://www.rplumber.io/>

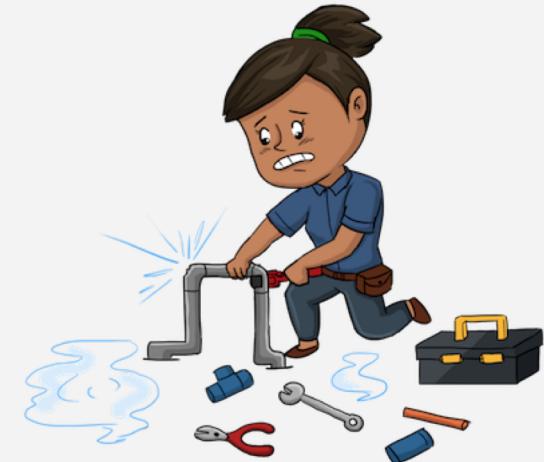
Building APIs with plumber

Step 2. Do scoring

score-lm.R

```
# read model
model <- readRDS("lm-model.rds")

## Score using linear model
## @param carat The carat value of the diamond
## @get /score-lm
function(carat) {
  list(price = predict(model,
    newdata = data.frame(
      carat = as.numeric(carat))))
}
```



<https://www.rplumber.io/>

Deploying

Base image

Default port 8000

Can customize

Build model in Dockerfile or
separately?

```
FROM trestletech/plumber
WORKDIR /R/
COPY . .
EXPOSE 8000

RUN ["install2.r", "ggplot2", "readr"]
RUN ["R", "-f", "/R/create-model.R"]

CMD [/R/score-lm.R]
```



```
$ docker build . -t score-lm

$ docker run -p 8000:8000 score-lm

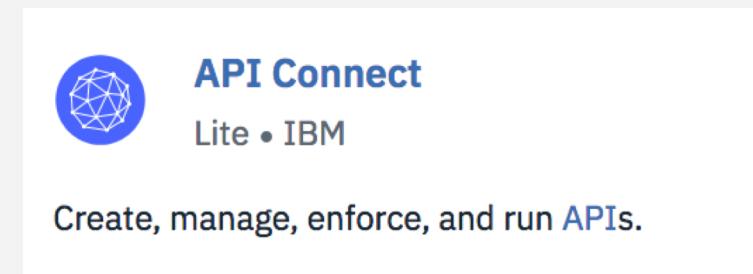
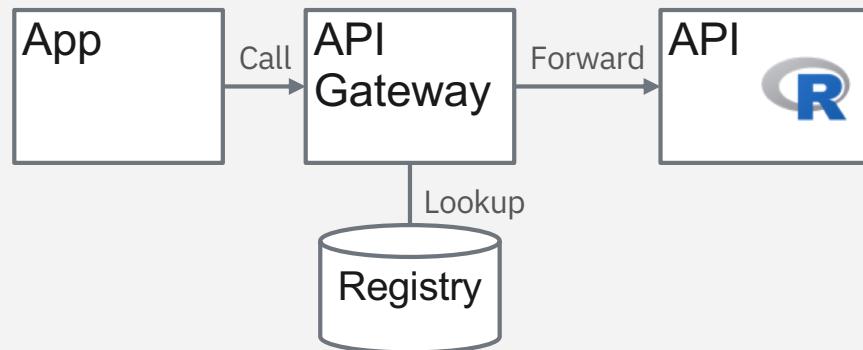
$ curl http://localhost:8000/score-lm?carat=0.5
{"price": [1621.8522]}
```

<https://hub.docker.com/r/trestletech/plumber/>

Securing APIs

1. Protect by network
2. Simple API key
3. API Gateway

- Considerations:
- Public-facing APIs
 - Brute force attacks
 - Denial of service attacks



Is R slow?

Compared to what?

What is too slow?

Several things we can do

Joe Cheng, CTO RStudio:

*“R can be slow, and is single-threaded –
But it’s probably less slow than you think”*

The image shows a screenshot of a ThoughtWorks blog post. The title is "Retail analytics: from hours to seconds using R". The author is Bharani Subramaniam, a market tech principal for India & Europe. The post is categorized under Technology, Data Science, and Data Science. A diagram illustrates a process flow: multiple "Week" stages (Week 1, Week 2, Week n) each produce an "R" (map), which then converge into a single "R" (reduce). Below the main image is a box containing the title and author information.

Timing the Same Algorithm in R, Python, and C++

While developing the `RcppDynProg` R package I took a little extra time to port the core algorithm from C++ to both R and Python.

<https://www.thoughtworks.com/insights/blog/retail-analytics-hours-seconds-using-r>
<http://www.win-vector.com/blog/2019/01/timing-of-the-same-algorithm-in-r-python-and-c/>

Write efficient code

Vectorize!

`data.table`

Profile to find
bottlenecks

`RCpp`

Leverage multi-core

`parallel`
`foreach`
`doParallel`

`ranger`

`caret`

`future / promises`

`ipc`

Scale with multiple
instances

Kubernetes –
load-balanced
pods

Node scaling

Map - Reduce

Integrate other
technologies



Resources

My tutorial: Deploying Shiny Apps on IBM Cloud

<https://github.com/holken1/deploying-r-on-cloud/tree/master/shiny-on-ibm-cloud>

- Free IBM Cloud cluster
- IBM Container Registry

My sample code:

<https://github.com/holken1/deploying-r-on-cloud>

IBM Cloud: <https://cloud.ibm.com>

RStudio: Load testing Shiny Applications

<https://rstudio.github.io/shinyloadtest/>

Joe Cheng: Shiny in Production (17 Jan)

<https://speakerdeck.com/jcheng5/shiny-in-production>

Colin Fay: An Introduction to Docker for R Users

<https://colinfay.me/docker-r-reproducibility/>

Yes, we can deploy R code.

It's easy!

Secure!

Scalable!



